

- Sahan Abeyrathna
- 210005H
- Assignment 1
- ENTC

Learning from data and related challenges and linear models for regression

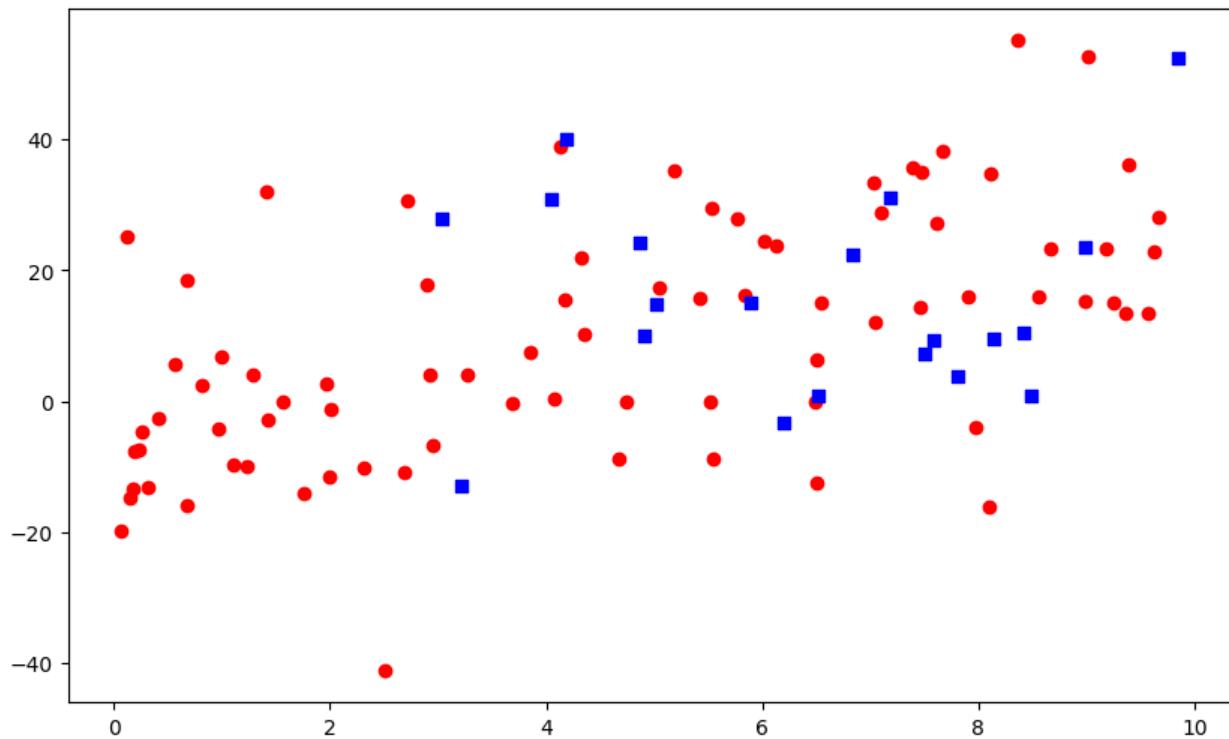
Learning from data

```

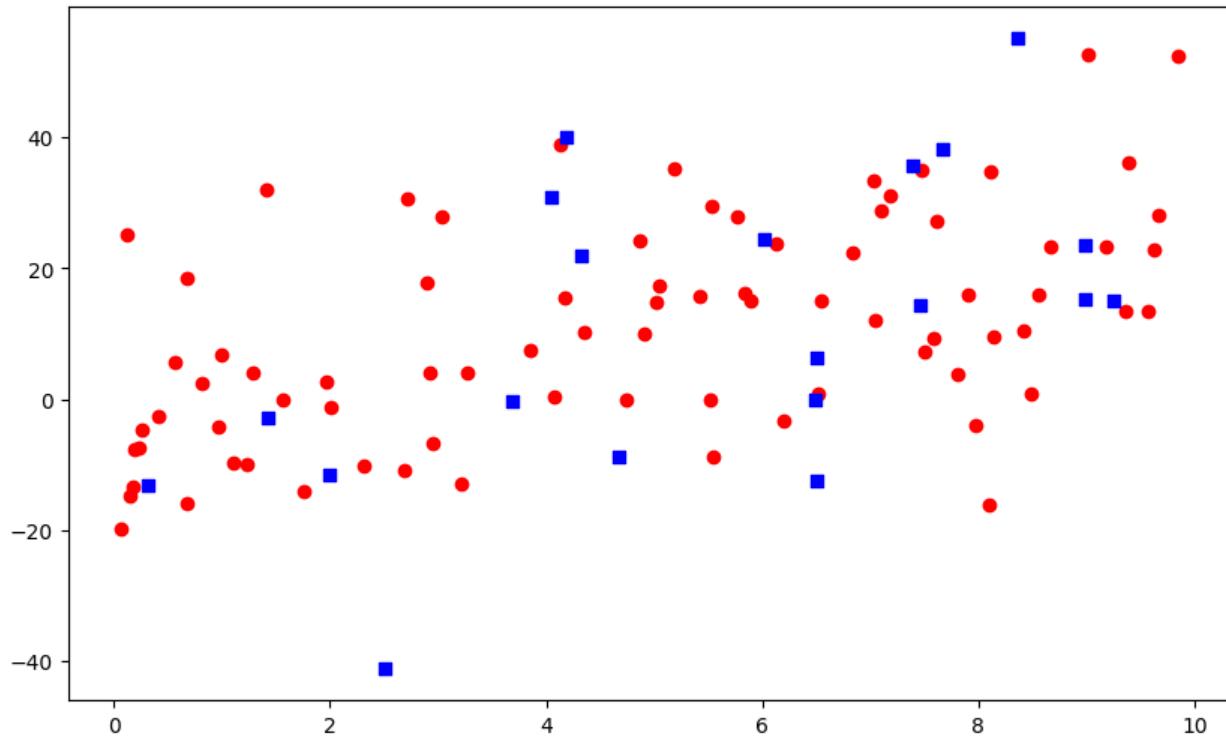
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
# Generate 100 samples
n_samples = 100
# Generate X values (uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)
# Generate epsilon values (normally distributed with mean 0 and standard deviation 15)
epsilon = np.random.normal(0, 15, n_samples)
# Generate Y values using the model Y = 3 + 3X + epsilon
Y = 3 + 2 * X + epsilon[:, np.newaxis]

r=np.random.randint(104)
# Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state=r)
# Plot the data points
plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1,
marker='o', color='red', label='Training Data')
plt.scatter(X_test, Y_test, alpha=1,
marker='s', color='blue', label='Testing Data')
plt.show()

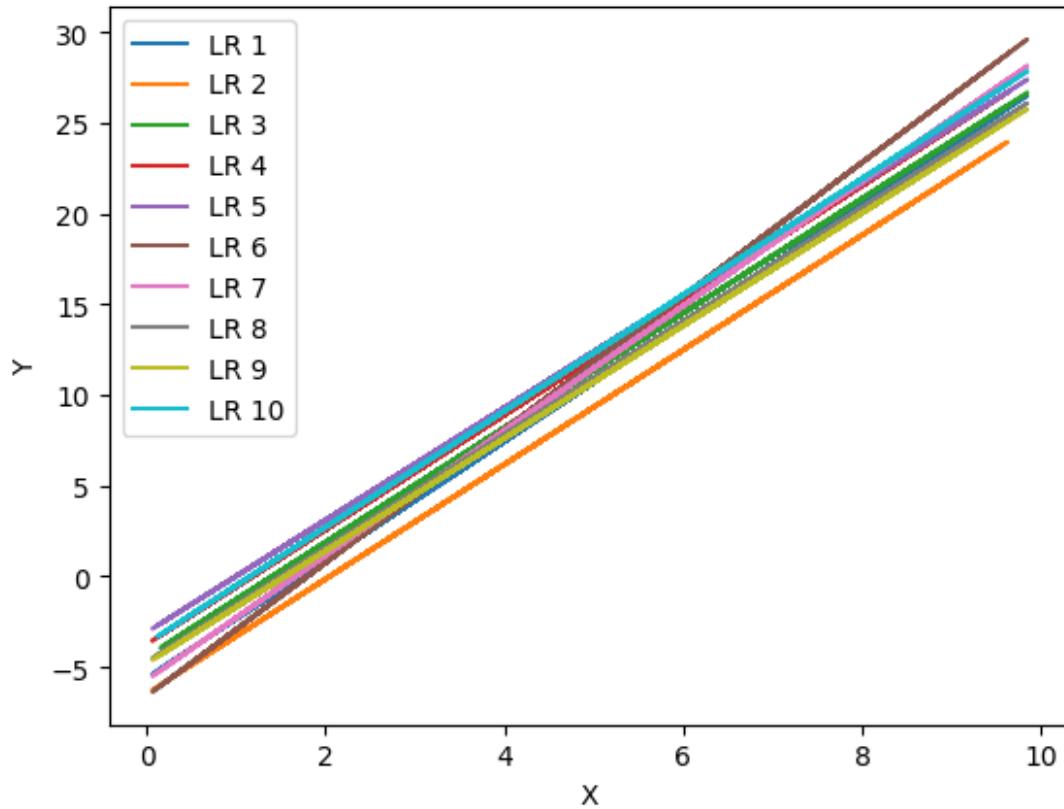
```



```
r=np.random.randint(104)
# Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X,
Y,test_size=0.2, random_state=r)
# Plot the data points
plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1,
marker='o',color='red',label='Training Data')
plt.scatter(X_test, Y_test, alpha=1,
marker='s',color='blue',label='Testing Data')
plt.show()
```



```
for i in range(10): # Plotting 10 different instances
    X_train, X_test, Y_train, Y_test = train_test_split(X,Y,
test_size=0.2, random_state=np.random.randint(104))
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```



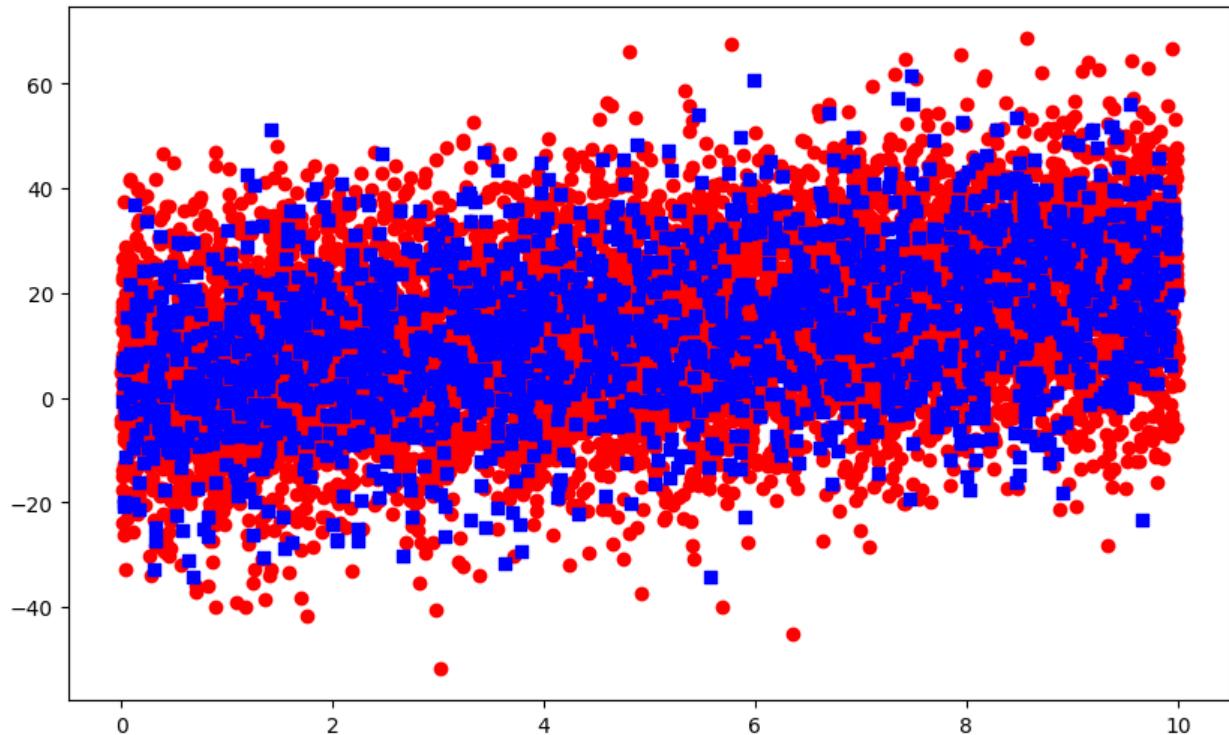
```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
# Generate 100 samples
n_samples = 10000
# Generate X values (uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)
# Generate epsilon values (normally distributed with mean 0 and
# standard deviation 15)
epsilon = np.random.normal(0, 15, n_samples)
# Generate Y values using the model Y = 3 + 3X + epsilon
Y = 3 + 2 * X + epsilon[:, np.newaxis]

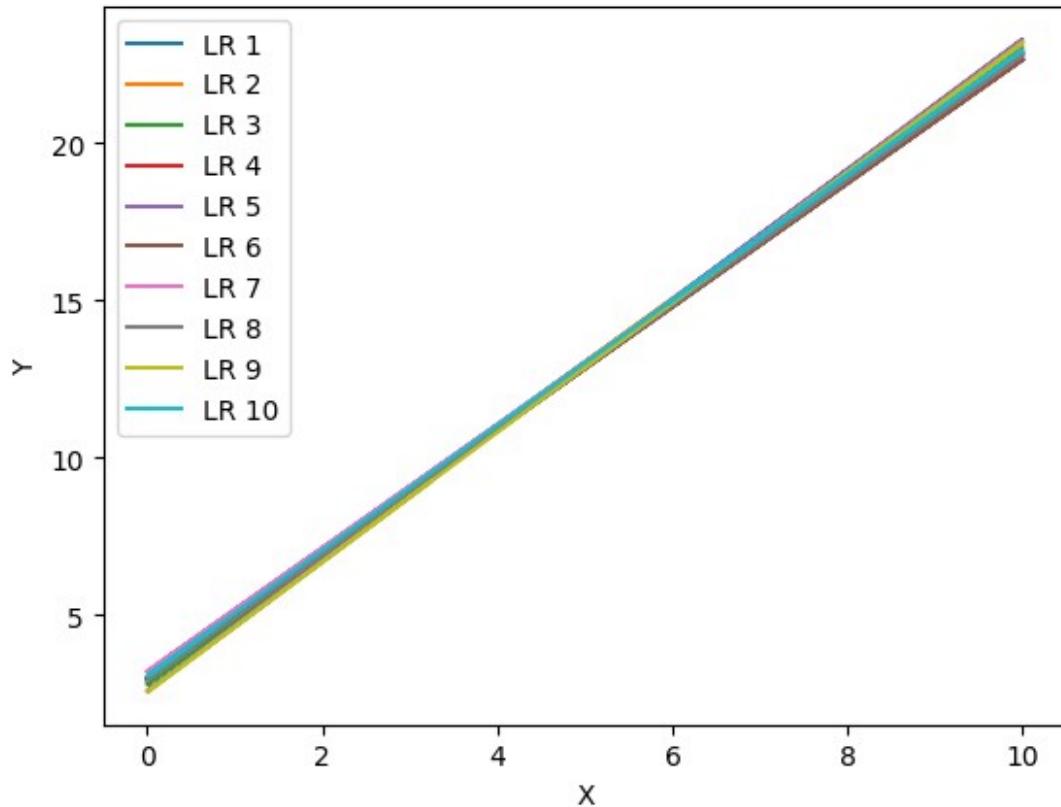
r=np.random.randint(104)
# Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state=r)
# Plot the data points
plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1,
marker='o', color='red', label='Training Data')
plt.scatter(X_test, Y_test, alpha=1,
marker='x', color='blue', label='Test Data')
# Plot the 10 linear regression models
for i in range(1, 11):
    lr = LinearRegression()
    lr.fit(X_train, Y_train)
    Y_pred = lr.predict(X_test)
    plt.plot(X_test, Y_pred, color=f'C{i}', label=f'LR {i}')
plt.legend()

```

```
marker='s',color='blue',label='Testing Data')
plt.show()
```



```
for i in range(10): # Plotting 10 different instances
    X_train, X_test, Y_train, Y_test = train_test_split(X,Y,
test_size=0.2, random_state=np.random.randint(104))
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```



```
pip install ucimlrepo
Collecting ucimlrepo
  Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: pandas>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2.1.4)
Requirement already satisfied: certifi>=2020.12.5 in
/usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2024.8.30)
Requirement already satisfied: numpy<2,>=1.22.4 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0-
>ucimlrepo) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0-
>ucimlrepo) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0-
>ucimlrepo) (2024.1)
Requirement already satisfied: tzdata>=2022.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0-
>ucimlrepo) (2024.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2-
>pandas>=1.0.0->ucimlrepo) (1.16.0)
Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
```

```
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7

from ucimlrepo import fetch_ucirepo
# fetch dataset
infrared_thermography_temperature = fetch_ucirepo(id=925)
# data (as pandas dataframes)
X = infrared_thermography_temperature.data.features
y = infrared_thermography_temperature.data.targets
# metadata
print(infrared_thermography_temperature.metadata)
# variable information
print(infrared_thermography_temperature.variables)

{'uci_id': 925, 'name': 'Infrared Thermography Temperature',
'repository_url':
'https://archive.ics.uci.edu/dataset/925/infrared+thermography+temperature+dataset', 'data_url':
'https://archive.ics.uci.edu/static/public/925/data.csv', 'abstract':
'The Infrared Thermography Temperature Dataset contains temperatures read from various locations of inferred images about patients, with the addition of oral temperatures measured for each individual. The 33 features consist of gender, age, ethnicity, ambient temperature, humidity, distance, and other temperature readings from the thermal images. The dataset is intended to be used in a regression task to predict the oral temperature using the environment information as well as the thermal image readings.', 'area': 'Health and Medicine',
'tasks': ['Regression'], 'characteristics': ['Tabular'],
'num_instances': 1020, 'num_features': 33, 'feature_types': ['Real',
'Categorical'], 'demographics': ['Gender', 'Age', 'Ethnicity'],
'target_col': ['ave0ralF', 'ave0ralM'], 'index_col': ['SubjectID'],
'has_missing_values': 'no', 'missing_values_symbol': None,
'year_of_dataset_creation': 2021, 'last_updated': 'Tue Dec 12 2023',
'dataset_doi': '10.13026/9ay4-2c37', 'creators': ['Quanzeng Wang',
'Yangling Zhou', 'Pejman Ghassemi', 'David McBride', 'J. Casamento',
'T. Pfefer', 'Quanzeng Wang', 'Yangling Zhou', 'Pejman Ghassemi',
'David McBride', 'J. Casamento', 'T. Pfefer'], 'intro_paper':
{'title': 'Infrared Thermography for Measuring Elevated Body Temperature: Clinical Accuracy, Calibration, and Evaluation',
'authors': 'Quanzeng Wang, Yangling Zhou, Pejman Ghassemi, David McBride, J. Casamento, T. Pfefer', 'published_in': 'Italian National Conference on Sensors', 'year': 2021, 'url':
'https://www.semanticscholar.org/paper/443b9932d295ca3a014e7d874b4bd77a33a276bd', 'doi': None}, 'additional_info': {'summary': None,
'purpose': None, 'funded_by': None, 'instances_represent': None,
'recommended_data_splits': None, 'sensitive_data': None,
'preprocessing_description': None, 'variable_info': '- gender\n- age\n- ethnicity\n- ambient temperature\n- humidity\n- distance\n- temperature readings from the thermal images', 'citation': None},
'external_url':
```


4	Age ranges in categories\n	None
no		
5	American Indian or Alaska Native, Asian, Black...	None
no		
6	Ambiant temperature	None
no		
7	Relative humidity	None
no		
8	Distance between the subjects and the IRTs.	None
no		
9	Temperature difference between the set and mea...	None
no		
10	Max value of a circle with diameter of 13 pixe...	None
no		
11	Max value of a circle with diameter of 13 pixe...	None
no		
12	Average value of a circle with diameter of 13 ...	None
no		
13	Average value of a circle with diameter of 13 ...	None
no		
14	Average temperature of the highest four pixels...	None
no		
15	Average temperature of the highest four pixels...	None
no		
16	Average temperature of the highest four pixels...	None
no		
17	Max value of a square of 24x24 pixels around t...	None
no		
18	Average temperature of the highest four pixels...	None
no		
19	Average temperature of the highest four pixels...	None
no		
20	Average temperature of the highest four pixels...	None
no		
21	Max value of a circle with diameter of 13 pixe...	None
no		
22	Average value of a square of 3x3 pixels center...	None
no		
23	Average value of a square of 3x3 pixels center...	None
no		
24	Max value in the extended canthi area	None
no		
25	Average temperature of the highest four pixels...	None
no		
26	Average value in the center point of forehead,...	None
no		
27	Average value in the right point of the forehe...	None
no		
28	Average value in the left point of the forehea...	None

```

no
29 Average value in the bottom point of the forehead. None
no
30 Average value in the top point of the forehead. None
no
31 Maximum temperature within the extended forehead. None
no
32 Max value in the center point of forehead, a s... None
no
33 Maximum temperature within the whole face region. None
no
34 Average temperature of the highest four pixels... None
no
35 Maximum temperature within the mouth region. None
no

print(infrared_thermography_temperature.data)

{'ids': SubjectID
0 161117-1
1 161117-2
2 161117-3
3 161117-4
4 161117-5
...
1015 180425-05
1016 180425-06
1017 180502-01
1018 180507-01
1019 180514-01

[1020 rows x 1 columns], 'features': Gender Age
Ethnicity T_atm Humidity Distance \
0 Male 41-50 White 24.0 28.0
0.8
1 Female 31-40 Black or African-American 24.0 26.0
0.8
2 Female 21-30 White 24.0 26.0
0.8
3 Female 21-30 Black or African-American 24.0 27.0
0.8
4 Male 18-20 White 24.0 27.0
0.8
...
...
1015 Female 21-25 Asian 25.7 50.8
0.6
1016 Female 21-25 White 25.7 50.8
0.6
1017 Female 18-20 Black or African-American 28.0 24.3

```

0.6							
1018	Male	26-30		Hispanic/Latino	25.0	39.8	
0.6							
1019	Female	18-20		White	23.8	45.6	
0.6							
T_FHRC1	T_offset1	Max1R13_1	Max1L13_1	aveAllR13_1	...	T_FHCC1	
0	\	0.7025	35.0300	35.3775	34.4000	...	33.5775
33.4775							
1	0.7800	34.5500	34.5200	33.9300	...	34.0325	
34.0550							
2	0.8625	35.6525	35.5175	34.2775	...	34.9000	
34.8275							
3	0.9300	35.2225	35.6125	34.3850	...	34.4400	
34.4225							
4	0.8950	35.5450	35.6650	34.9100	...	35.0900	
35.1600							
...
1015	1.2225	35.6425	35.6525	34.8575	...	35.1075	
35.3475							
1016	1.4675	35.9825	35.7575	35.4275	...	35.3100	
35.2175							
1017	0.1300	36.4075	36.3400	35.8700	...	35.4350	
35.2400							
1018	1.2450	35.8150	35.5250	34.2950	...	34.8400	
35.0200							
1019	0.8675	35.7075	35.5825	34.8875	...	34.5475	
34.6500							
T_OR1	T_FHLC1	T_FHBC1	T_FHTC1	T_FH_Max1	T_FHC_Max1	T_Max1	
0	\	33.3725	33.4925	33.0025	34.5300	34.0075	35.6925
35.6350							
1	33.6775	33.9700	34.0025	34.6825	34.6600	35.1750	
35.0925							
2	34.6475	34.8200	34.6700	35.3450	35.2225	35.9125	
35.8600							
3	34.6550	34.3025	34.9175	35.6025	35.3150	35.7200	
34.9650							
4	34.3975	34.6700	33.8275	35.4175	35.3725	35.8950	
35.5875							
...
1015	35.4000	35.1375	35.2750	35.8525	35.7475	36.0675	
35.6775							
1016	35.2200	35.2075	35.0700	35.7650	35.5525	36.5000	
36.4525							

```
1017 35.2275 35.3675 35.3425      36.3750      35.7100 36.5350  
35.9650  
1018 34.9250 34.7150 34.5950      35.4150      35.3100 35.8600  
35.4150  
1019 34.6700 34.2150 34.7100      35.1525      35.1175 35.9725  
35.8900
```

```
T_OR_Max1  
0      35.6525  
1      35.1075  
2      35.8850  
3      34.9825  
4      35.6175  
...  
1015 35.7100  
1016 36.4900  
1017 35.9975  
1018 35.4350  
1019 35.9175
```

```
[1020 rows x 33 columns], 'targets':      ave0ralF  ave0ralM  
0      36.85      36.59  
1      37.00      37.19  
2      37.20      37.34  
3      36.85      37.09  
4      36.80      37.04  
...  
1015 36.95      36.99  
1016 37.25      37.19  
1017 37.35      37.59  
1018 37.15      37.29  
1019 37.05      37.19
```

```
[1020 rows x 2 columns], 'original':      SubjectID  ave0ralF  
ave0ralM  Gender    Age          Ethnicity \\\n0      161117-1  36.85      36.59      Male   41-50  
White  
1      161117-2  37.00      37.19      Female  31-40  Black or African-  
American  
2      161117-3  37.20      37.34      Female  21-30  
White  
3      161117-4  36.85      37.09      Female  21-30  Black or African-  
American  
4      161117-5  36.80      37.04      Male   18-20  
White  
...  
1015 180425-05  36.95      36.99      Female  21-25  
Asian  
1016 180425-06  37.25      37.19      Female  21-25
```



```

1016 35.2075 35.0700    35.7650    35.5525 36.5000 36.4525
36.4900
1017 35.3675 35.3425    36.3750    35.7100 36.5350 35.9650
35.9975
1018 34.7150 34.5950    35.4150    35.3100 35.8600 35.4150
35.4350
1019 34.2150 34.7100    35.1525    35.1175 35.9725 35.8900
35.9175

[1020 rows x 36 columns], 'headers': Index(['SubjectID', 'aveOralF',
'aveOralM', 'Gender', 'Age', 'Ethnicity',
'T_atm', 'Humidity', 'Distance', 'T_offset1', 'Max1R13_1',
'Max1L13_1',
'aveAllR13_1', 'aveAllL13_1', 'T_RC1', 'T_RC_Dry1',
'T_RC_Wet1',
'T_RC_Max1', 'T_LC1', 'T_LC_Dry1', 'T_LC_Wet1', 'T_LC_Max1',
'RCC1',
'LCC1', 'canthiMax1', 'canthi4Max1', 'T_FHCC1', 'T_FHRC1',
'T_FHLC1',
'T_FHBC1', 'T_FHTC1', 'T_FH_Max1', 'T_FHC_Max1', 'T_Max1',
'T_OR1',
'T_OR_Max1'],
dtype='object')}

import pandas as pd
print(f"X shape beforeremoval: {X.shape}")
print(f"y shape beforeremoval: {y.shape}")
df =pd.concat([X,y],axis =1)
df =df.dropna()
X =df.iloc[:, :-2]
y =df.iloc[:, -2:]
print(f"X shape afterremoval:{X.shape}")
print(f"y shape afterremoval:{y.shape}")

X shape beforeremoval: (1020, 33)
y shape beforeremoval: (1020, 2)
X shape afterremoval:(1018, 33)
y shape afterremoval:(1018, 2)

target =y['aveOralM']
features= X[['Age', 'T_atm','Humidity','T_Max1','T_offset1']]
print(features)

      Age   T_atm  Humidity   T_Max1   T_offset1
0    41-50    24.0     28.0    35.6925    0.7025
1    31-40    24.0     26.0    35.1750    0.7800
2    21-30    24.0     26.0    35.9125    0.8625
3    21-30    24.0     27.0    35.7200    0.9300
4    18-20    24.0     27.0    35.8950    0.8950
...

```

```

1015 21-25 25.7      50.8 36.0675    1.2225
1016 21-25 25.7      50.8 36.5000    1.4675
1017 18-20 28.0      24.3 36.5350    0.1300
1018 26-30 25.0      39.8 35.8600    1.2450
1019 18-20 23.8      45.6 35.9725    0.8675

```

[1018 rows x 5 columns]

```

def convert_age_range(age_range):
    """Converts the age range to a single average value"""
    if '>' in age_range:
        return int(age_range.replace('>','')).strip()
    lower,upper =map(int,age_range.split('-'))
    return np.mean([lower,upper])

```

```

features['Age']= features['Age'].apply(convert_age_range)
print(features)

```

	Age	T_atm	Humidity	T_Max1	T_offset1
0	45.5	24.0	28.0	35.6925	0.7025
1	35.5	24.0	26.0	35.1750	0.7800
2	25.5	24.0	26.0	35.9125	0.8625
3	25.5	24.0	27.0	35.7200	0.9300
4	19.0	24.0	27.0	35.8950	0.8950
..
1015	23.0	25.7	50.8	36.0675	1.2225
1016	23.0	25.7	50.8	36.5000	1.4675
1017	19.0	28.0	24.3	36.5350	0.1300
1018	28.0	25.0	39.8	35.8600	1.2450
1019	19.0	23.8	45.6	35.9725	0.8675

[1018 rows x 5 columns]

```

<ipython-input-23-1a38761dabf6>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

features['Age']= features['Age'].apply(convert_age_range)

```

```

X_train,X_test, y_train,y_test=
train_test_split(features,target,test_size =
0.2,random_state=np.random.randint(104))

```

```

linear_regression =LinearRegression()
linear_regression.fit(X_train,y_train)
print('Intercept:\n',linear_regression.intercept_)
print('Coefficients: \n',linear_regression.coef_)

```

```

Intercept:
5.350271904448654
Coefficients:
 [ 0.00106708 -0.06026008  0.00099853  0.91390548  0.09962315]

features= X[['T_OR1','T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']]
X_train,X_test, y_train,y_test=
train_test_split(features,target,test_size =
0.2,random_state=np.random.randint(104))
linear_regression =LinearRegression()
linear_regression.fit(X_train,y_train)
print('Intercept:\n',linear_regression.intercept_)
print('Coefficients: \n',linear_regression.coef_)

Intercept:
7.1436575396059006
Coefficients:
 [-0.05575003  0.60320041 -0.05115736  0.34035037]

print(X_test.shape)

(204, 4)

import statsmodels.api as sm

residual_sum_of_squares = np.sum((linear_regression.predict(X_test)-
y_test)** 2)
residual_standard_error = np.sqrt(residual_sum_of_squares /
(X_test.shape[0]-X_test.shape[1]- 1))

mean_squared_error = np.mean((linear_regression.predict(X_test)-
y_test) ** 2)

r_squared = linear_regression.score(X_test, y_test)
X_test_with_intercept = np.c_[np.ones(X_test.shape[0]), X_test] # Add
intercept term if not included
XtX_inv = np.linalg.inv(X_test_with_intercept.T @
X_test_with_intercept)
standard_error_coefficients = np.sqrt(np.diag(residual_standard_error
** 2 * XtX_inv))

ols_model = sm.OLS(y_test, X_test_with_intercept).fit()
summary = ols_model.summary()

t_values = ols_model.tvalues.values
p_values = ols_model.pvalues.values

print(f"Residual Sum of Squares: {residual_sum_of_squares}\n")
print(f"Residual Standard Error: {residual_standard_error}\n")
print(f"Mean Squared Error: {mean_squared_error}\n")
print(f"R-Squared: {r_squared}\n")

```

```
print(f"Standard Error of Coefficients:\n{standard_error_coefficients}\n")
print(f"t-values:\n{t_values}\n")
print(f"p-values:\n{p_values}\n")
print(f"\n\n\n\n{summary}")

Residual Sum of Squares: 20.30989897585181
Residual Standard Error: 0.31946798563940887
Mean Squared Error: 0.09955832831299906
R-Squared: 0.6428739418600805
Standard Error of Coefficients:
[1.56946255 1.80438807 1.79706711 0.08663473 0.09491936]

t-values:
[ 5.51634926  0.76840282 -0.46550164 -1.15883401  3.70343768]

p-values:
[1.06823189e-07 4.43158815e-01 6.42081296e-01 2.47912556e-01
 2.75451148e-04]
```

OLS Regression Results

```
=====
=====
Dep. Variable: ave0ralM R-squared: 0.647
Model: OLS Adj. R-squared: 0.640
Method: Least Squares F-statistic: 91.15
Date: Tue, 10 Sep 2024 Prob (F-statistic): 6.78e-44
Time: 13:47:16 Log-Likelihood: -52.991
No. Observations: 204 AIC: 116.0
Df Residuals: 199 BIC: 132.6
Df Model: 4
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025
0.975]					
const	8.6087	1.561	5.516	0.000	5.531
11.686					
x1	1.3787	1.794	0.768	0.443	-2.159
4.917					
x2	-0.8318	1.787	-0.466	0.642	-4.356
2.692					
x3	-0.0998	0.086	-1.159	0.248	-0.270
0.070					
x4	0.3495	0.094	3.703	0.000	0.163
0.536					
=====					
Omnibus:		25.345	Durbin-Watson:		
1.857					
Prob(Omnibus):		0.000	Jarque-Bera (JB):		
41.636					
Skew:		0.684	Prob(JB):		
9.09e-10					
Kurtosis:		4.740	Cond. No.		
8.11e+03					
=====					
=====					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.11e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
r =np.linspace(-100, 100,400)
#Definethe lossfunctions
def L1(r,a):
    return r**2/ (a**2 +r**2)

def L2(r,a):
    return 1-np.exp(-2 *np.abs(r) / a)

# Define different values of a

a_values = [2, 5, 10, 20, 30, 50, 100]

# Plotting

plt.figure(figsize=(14, 6))
```

```

# Plot L1(r) for different values of a
plt.subplot(1, 2, 1)
for a in a_values:
    plt.plot(r, L1(r, a), label=f'a = {a}')
plt.title('$L_1(r)$')
plt.xlabel('Residual $r$')
plt.ylabel('$L_1(r)$')
plt.legend()
plt.grid(True)
# Plot L2(r) for different values of a
plt.subplot(1, 2, 2)
for a in a_values:
    plt.plot(r, L2(r, a), label=f'a = {a}')
plt.title('$L_2(r)$')
plt.xlabel('Residual $r$')
plt.ylabel('$L_2(r)$')
plt.legend()
plt.grid(True)

```

